


Importing the Dependencies

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

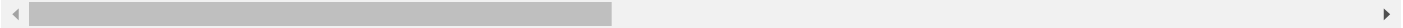
# loading the dataset to a Pandas DataFrame
credit_card_data = pd.read_csv('/content/creditcard.csv')

# First 5 rows of the dataset
credit_card_data.head()
```




	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23
0	0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277838	-0.110474
1	0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288
2	1	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679	0.909412
3	1	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321
4	2	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458

5 rows × 31 columns

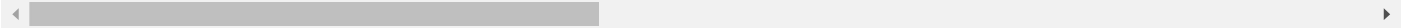


```
credit_card_data.tail()
```




	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V2
29794	35633	0.786689	-0.691214	-0.329291	0.149435	0.714779	1.949061	-0.136906	0.474172	0.206173	...	-0.165285	-0.793473	-0.03011
29795	35633	0.800996	-2.159993	0.008378	-1.081828	-1.768799	-0.445016	-0.571165	-0.162429	-1.785636	...	0.016930	-0.350492	-0.23488
29796	35633	1.115726	-0.472602	0.983034	0.294673	-1.218768	-0.341755	-0.667340	0.171155	0.805427	...	0.104463	0.366801	-0.07321
29797	35634	1.239103	-1.000617	0.843324	-0.560021	-1.400343	-0.151696	-1.026058	-0.001637	-0.131138	...	0.325954	0.855203	-0.24568
29798	35634	1.374193	-0.720679	0.891375	-0.541402	-1.700000	NaN	NaN	NaN	NaN	...	NaN	NaN	NaN

5 rows × 31 columns



```
#dataset informations
credit_card_data.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 29799 entries, 0 to 29798
Data columns (total 31 columns):
#   Column      Non-Null Count  Dtype
---  ---
0    Time        29799 non-null  int64
1    V1           29799 non-null  float64
2    V2           29799 non-null  float64
3    V3           29799 non-null  float64
4    V4           29799 non-null  float64
5    V5           29799 non-null  float64
6    V6           29798 non-null  float64
7    V7           29798 non-null  float64
8    V8           29798 non-null  float64
9    V9           29798 non-null  float64
10   V10          29798 non-null  float64
11   V11          29798 non-null  float64
12   V12          29798 non-null  float64
13   V13          29798 non-null  float64
14   V14          29798 non-null  float64
15   V15          29798 non-null  float64
16   V16          29798 non-null  float64
17   V17          29798 non-null  float64
18   V18          29798 non-null  float64
19   V19          29798 non-null  float64
20   V20          29798 non-null  float64
```

```

21 V21      29798 non-null float64
22 V22      29798 non-null float64
23 V23      29798 non-null float64
24 V24      29798 non-null float64
25 V25      29798 non-null float64
26 V26      29798 non-null float64
27 V27      29798 non-null float64
28 V28      29798 non-null float64
29 Amount    29798 non-null float64
30 Class     29798 non-null float64
dtypes: float64(30), int64(1)
memory usage: 7.0 MB

```

```

#checking the number of missing values in each column
credit_card_data.isnull().sum()

```

```

Time      0
V1        0
V2        0
V3        0
V4        0
V5        0
V6        1
V7        1
V8        1
V9        1
V10       1
V11       1
V12       1
V13       1
V14       1
V15       1
V16       1
V17       1
V18       1
V19       1
V20       1
V21       1
V22       1
V23       1
V24       1
V25       1
V26       1
V27       1
V28       1
Amount    1
Class     1
dtype: int64

```

```

#distribution of legit transactions & fraudulent transactions
credit_card_data['Class'].value_counts()

```

```

Class
0.0    29704
1.0      94
Name: count, dtype: int64

```

This Dataset is highly unbalanced

0--> Normal Transaction

1--> Fraudulent Transaction

```

#separating the data for analysis

```

```

legit = credit_card_data[credit_card_data.Class == 0]
fraud = credit_card_data[credit_card_data.Class == 1]

```

```

print(legit.shape)
print(fraud.shape)

```

```

(29704, 31)
(94, 31)

```

```

#statistical measures of the data
legit.Amount.describe()

```

```
count    29704.000000
mean      79.570030
std       221.991154
min        0.000000
25%        6.637500
50%       20.000000
75%       70.652500
max      7879.420000
Name: Amount, dtype: float64
```

```
fraud.Amount.describe()
```

```
count     94.000000
mean      95.590000
std       257.920621
min        0.000000
25%        1.000000
50%       1.050000
75%       99.990000
max     1809.680000
Name: Amount, dtype: float64
```

```
#compare the values for both transactions
credit_card_data.groupby('Class').mean()
```

```

      Time      V1      V2      V3      V4      V5      V6      V7      V8      V9  ...  V20      V21
Class
0.0  21422.566422 -0.184395  0.106639  0.759413  0.194690 -0.186422  0.096790 -0.096841  0.018203  0.361273  ...  0.044000 -0.035795 -0.035795
1.0   19007.702128 -8.099702  6.084984 -11.565958  6.014185 -5.681925 -2.370349 -7.912202  4.043743 -2.891421  ...  0.679513  0.573983 -0.573983
2 rows x 30 columns
```

Under-Sampling

Build a sample dataset containing similar distribution of normal transactions and Fraudulent Transactions

Number of Fraudulent Transactions --> 492

```
legit_sample = legit.sample(n=492)
```

Concatenating two DataFrames


```
new_dataset = pd.concat([legit_sample, fraud], axis=0)
```

```
new_dataset.head()
```

```

      Time      V1      V2      V3      V4      V5      V6      V7      V8      V9  ...  V21      V22
7755  10799  1.266542  0.015942  0.623743 -0.050315 -0.513566 -0.469132 -0.423446 -0.067397  1.470987  ... -0.256529 -0.522403  0.056
22921  32503 -1.950713  1.976018  0.718646 -0.084854 -0.389506 -0.873640  0.166806 -1.740505  0.081659  ...  1.628822 -0.023158  0.179
16292  27689  0.808134 -0.364113  0.106022  1.335512 -0.045837  0.410419  0.231907  0.072522 -0.060674  ...  0.041210 -0.045832 -0.276
12943  22744 -11.050688  6.768340 -11.068786  2.497231 -6.750103 -2.797250 -4.525065  7.175295  1.199926  ... -0.171703 -1.230148 -0.376
7864  10945  1.261789 -0.058362  0.552553  0.015400 -0.358434 -0.039203 -0.544068 -0.014109  1.506255  ... -0.067001  0.044749 -0.136
5 rows x 31 columns
```


```
new_dataset.tail()
```



	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23
27362	34521	1.081234	0.416414	0.862919	2.520863	-0.005021	0.563341	-0.123372	0.223122	-0.673598	...	-0.159387	-0.305154	0.053621
27627	34634	0.333499	1.699873	-2.596561	3.643945	-0.585068	-0.654659	-2.275789	0.675229	-2.042416	...	0.469212	-0.144363	-0.317987
27738	34684	-2.439237	2.591458	-2.840126	1.286244	-1.777016	-1.436139	-2.206056	-2.282725	-0.292885	...	1.774460	-0.771390	0.065727
27749	34687	-0.860827	3.131790	-5.052968	5.420941	-2.494141	-1.811287	-5.479117	1.189472	-3.908206	...	1.192694	0.090356	-0.341887
29687	35585	-2.019001	1.491270	0.005222	0.817253	0.973252	-0.639268	-0.974073	-3.146929	-0.003159	...	2.839596	-1.185443	-0.142817


5 rows × 31 columns

```
new_dataset['Class'].value_counts()
```



Class	
0.0	492
1.0	94
Name: count, dtype: int64	

```
new_dataset.groupby('Class').mean()
```




	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V20	V21	V22
Class														
0.0	21659.467480	-0.126652	0.159711	0.648066	0.132789	-0.223356	0.054391	-0.116954	0.051761	0.285043	...	0.079924	-0.040564	-0.000000
1.0	19007.702128	-8.099702	6.084984	-11.565958	6.014185	-5.681925	-2.370349	-7.912202	4.043743	-2.891421	...	0.679513	0.573983	-0.000000

2 rows × 30 columns

Splitting the data into Features & Targets

```
X = new_dataset.drop(columns='Class', axis=1)
Y = new_dataset['Class']
```

```
print(X)
```



	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V20	V21	V22	V23	V24	V25	V26	V27	V28	Amount
7755	10799	1.266542	0.015942	0.623743	-0.050315	-0.513566	-0.469132														
22921	32503	-1.950713	1.976018	0.718646	-0.084854	-0.389506	-0.873640														
16292	27689	0.808134	-0.364113	0.106022	1.335512	-0.045837	0.410419														
12943	22744	-11.050688	6.768340	-11.068786	2.497231	-6.750103	-2.797250														
7864	10945	1.261789	-0.058362	0.552553	0.015400	-0.358434	-0.039203														
...														
27362	34521	1.081234	0.416414	0.862919	2.520863	-0.005021	0.563341														
27627	34634	0.333499	1.699873	-2.596561	3.643945	-0.585068	-0.654659														
27738	34684	-2.439237	2.591458	-2.840126	1.286244	-1.777016	-1.436139														
27749	34687	-0.860827	3.131790	-5.052968	5.420941	-2.494141	-1.811287														
29687	35585	-2.019001	1.491270	0.005222	0.817253	0.973252	-0.639268														
7755																					
22921																					
16292																					
12943																					
7864																					
...																					
27362																					
27627																					
27738																					
27749																					
29687																					
7755																					
22921																					
16292																					
12943																					
7864																					
...																					
27362																					
27627																					
27738																					
27749																					
29687																					

```
29687 -0.142812 -0.086103 -0.329113 0.523601 0.626283 0.152440 0.76
```

```
[586 rows x 30 columns]
```

```
print(Y)
```

```
7755    0.0
22921    0.0
16292    0.0
12943    0.0
7864     0.0
...
27362    1.0
27627    1.0
27738    1.0
27749    1.0
29687    1.0
Name: Class, Length: 586, dtype: float64
```

Split the data into Training data & Testing Data

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, stratify=Y, random_state=2)
```

```
print(X.shape, X_train.shape, X_test.shape)
```

```
(586, 30) (468, 30) (118, 30)
```

Model Training

Logistic Regression

```
model = LogisticRegression()
```

```
# training the Logistic Regression Model with Training Data
```

```
model.fit(X_train, Y_train)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
    LogisticRegression
```

Model Evaluation

Accuracy Score

```
# accuracy on training data
```

```
X_train_prediction = model.predict(X_train)
```

```
training_data_accuracy = accuracy_score(X_train_prediction, Y_train)
```

```
print('Accuracy on Training data : ', training_data_accuracy)
```

```
Accuracy on Training data : 0.9722222222222222
```

```
# accuracy on test data
```

```
X_test_prediction = model.predict(X_test)
```

```
test_data_accuracy = accuracy_score(X_test_prediction, Y_test)
```

```
print('Accuracy score on Test Data : ', test_data_accuracy)
```

```
Accuracy score on Test Data : 0.9322033898305084
```

